

Outlook

From	david.fourie@infrastructure.ops.keystonebank.co.za
To	analytics.retail@keystonebank.co.za, banking.operations@keystonebank.co.za
Cc	customer.experience@keystonebank.co.za, cyber.security@keystonebank.co.za
Sent	Tue, 22 Oct 2019 10:22:00 -0000

The incident clock and the customer clock are not the same

Hi,

This came up in a working session and nobody had the same definition.

Infrastructure can see technical symptoms, but incident severity should reflect failed customer journeys and repeat attempts rather than machine alerts alone.

I do not want a dashboard that simply counts the outcome. Start with the competing explanations: latency rises before visible failure; customers retry across channels and create duplicates; or a small terminal or channel group creates most harm.

Please start with the decision we need to make, then show whether the available evidence is strong enough to support it.

What we need to decide is incident priority, root-cause ownership and the customer-remediation population. Please bring a working Excel file with the exception population and clear columns, including a few cases we can trace from source to conclusion.

The available evidence is transaction-level timestamps, channels, amounts, statuses and error context; ATM attempts, host responses, PIN attempts, blocked-card state and latency; customer contacts, complaints and suggestions. The dataset has channel and ATM telemetry, not complete application, server or network logs.

Thank you