

Outlook

From	david.fourie@infrastructure.ops.keystonebank.co.za
To	analytics.retail@keystonebank.co.za, banking.operations@keystonebank.co.za
Cc	customer.experience@keystonebank.co.za, cyber.security@keystonebank.co.za
Sent	Tue, 05 Oct 2021 09:22:00 -0000

Follow-up: the incident clock and the customer clock are not the same

Morning,

I need help with something that sounds simple but probably is not.

Infrastructure can see technical symptoms, but incident severity should reflect failed customer journeys and repeat attempts rather than machine alerts alone. Several teams are reporting more failed or retried activity than they expected, but they do not agree on the cause. Different teams have produced different answers from the same period, so population and timing rules need to be visible.

The questions I would ask in the room are:

1. What would we expect to see if customers retry across channels and create duplicates?
2. What evidence would instead support that latency rises before visible failure?
3. How do we rule out that a small terminal or channel group creates most harm?

This has returned because the previous answer was useful but not strong enough to become a standing rule. Use October 2021 as the new evidence point rather than recycling the earlier conclusion. Please work towards incident priority, root-cause ownership and the customer-remediation population. A useful output would be a model-ready table, data dictionary and an honest baseline.

Use transaction-level timestamps, channels, amounts, statuses and error context; ATM attempts, host responses, PIN attempts, blocked-card state and latency; customer contacts, complaints and suggestions. One caution: The dataset has channel and ATM telemetry, not complete application, server or network logs.

Regards